

# Database Management Systems - I, CS 157A

**Transactions in SQL &  
Constraints and Triggers  
& Views and Indexes**

**Ch. 6.6, Ch. 7, Ch. 8**



# Agenda

- Transactions Overview
- Constraints Overview
- Triggers Overview
- Views Overview
- Index Overview



# **Transactions Overview**

## **Ch. 6.6**

# Why Transactions?

- Database systems are normally being accessed by many users or processes at the same time:
  - Both queries and modifications
- Unlike operating systems, which *support* interaction of processes, a DMBS needs to keep processes from troublesome interactions

# Example: Bad Interaction

- You and your domestic partner each take \$100 from different ATM's at about the same time:
  - The DBMS better make sure one account deduction doesn't get lost
- **Compare:** An OS allows two people to edit a document at the same time. If both write, one's changes get lost.

# Transactions

- ***Transaction*** = set of operations as a unit of work involving database queries and/or modification
- Normally with some strong properties regarding concurrency
- Formed in SQL from single statements or explicit programmer control (transaction demarcation)

# ACID Transactions

## ■ *ACID transactions are:*

- *Atomic*: Whole transaction or none is done
- *Consistent*: Database constraints preserved
- *Isolated*: It appears to the user as if only one process executes at a time
- *Durable*: Effects of a process survive a crash

## ■ **Optional:** weaker forms of transactions are often supported as well

# COMMIT

- The SQL statement COMMIT causes a transaction to complete:
  - It's database modifications are now permanent in the database



# ROLLBACK

- The SQL statement **ROLLBACK** also causes the transaction to end, but by *aborting*:
  - No effects on the database
- Failures like division by 0 or a constraint violation can also cause rollback, even if the programmer does not request it

# Example: Interacting Processes

- Assume the usual **Sells**(bar,beer,price) relation, and suppose that Joe's Bar sells only Bud for \$2.50 and Miller for \$3.00
- Sally is querying **Sells** for the highest and lowest price Joe charges
- Joe decides to stop selling Bud and Miller, but to sell only Heineken at \$3.50

# Sally's Program

- Sally executes the following two SQL statements called (min) and (max) to help us remember what they do:

(max)      **SELECT** MAX(price) **FROM** Sells  
             **WHERE** bar = 'Joe's Bar';

(min)      **SELECT** MIN(price) **FROM** Sells  
             **WHERE** bar = 'Joe's Bar';

# Joe's Program

- At about the same time, Joe executes the following steps: (del) and (ins):

(del) **DELETE FROM Sells**  
**WHERE** bar = 'Joe''s Bar';

(ins) **INSERT INTO Sells**  
**VALUES**('Joe''s Bar', 'Heineken', 3.50);

# Interleaving of Statements

- Although (max) must come before (min), and (del) must come before (ins), there are no other constraints on the order of these statements, unless we group Sally's and/or Joe's statements into transactions

# Example: Strange Interleaving

- Suppose the steps execute in the order  
(max)(del)(ins)(min):

Joe's Prices:

Statement	{2.50,3.00}	{2.50,3.00}		{3.50}
Result:	(max)	(del)	(ins)	(min)
	3.00			3.50

- Sally sees MAX < MIN!

# Fixing the Problem by Using Transactions

- If we group Sally's statements (max)(min) into one transaction, then she will not see this inconsistency
- She sees Joe's prices at some fixed time:
  - Either before or after he changes prices, or in the middle, but the MAX and MIN are computed from the same prices

# Another Problem: Rollback

- Suppose Joe executes (del)(ins), not as a transaction, but after executing these statements, thinks better of it and issues a ROLLBACK statement
- If Sally executes her statements after (ins) but before the rollback, she sees a value, 3.50, that never existed in the database because of the rollback



# Solution

- If Joe executes (del)(ins) as a transaction, its effect cannot be seen by others until the transaction executes COMMIT:
  - If the transaction executes ROLLBACK instead, then its effects will *never* be seen

# Isolation Levels

- SQL defines four *isolation levels* = choices about what interactions are allowed by transactions that execute at about the same time
- Only one level (“*serializable*”) = ACID transactions
- Each DBMS implements transactions in its own way

# Choosing the Isolation Level

- Within a transaction, we can say:  
**SET TRANSACTION ISOLATION  
LEVEL *X***

where *X* =

1. SERIALIZABLE
2. REPEATABLE READ
3. READ COMMITTED
4. READ UNCOMMITTED

# Serializable Transactions

- If Sally = (max)(min) and Joe = (del)(ins) are each transactions, and Sally runs with isolation level **SERIALIZABLE**, then she will see the database either before or after Joe runs, but not in the middle

# Isolation Level Is Personal Choice

- Your choice, e.g., run **serializable**, affects only how *you* see the database, not how others see it
- **Example**: If **Joe** Runs **serializable**, but **Sally** doesn't, then **Sally** might see no prices for Joe's Bar:
  - i.e., it looks to **Sally** as if she ran in the middle of **Joe**'s transaction, i.e. after delete but before insert

# Read-Committed Transactions

- If Sally runs with isolation level **READ COMMITTED**, then she can see only committed data, but not necessarily the same data each time
- **Example:** Under **READ COMMITTED**, the interleaving **(max)(del)(ins)(min)** is allowed, as long as Joe commits:
  - sees  $MAX < MIN$

# Repeatable-Read Transactions

- Requirement is like **read-committed**, plus.  
if data is read again, then everything seen  
the first time will be seen the second time:
  - But the second and subsequent reads may  
see *more* tuples as well

# Example: Repeatable Read

- Suppose Sally (Max, Min) runs under **REPEATABLE READ**, and the order of execution is **(max)(del)(ins)(min)**:
  - **(max)** sees prices 2.50 and 3.00
  - **(min)** can see 3.50, but must also see 2.50 and 3.00, because they were seen on the earlier read by **(max)**



# Read Uncommitted

- A transaction running under **READ UNCOMMITTED** can see data in the database, even if it was written by a transaction that has not committed (and may never)
- **Example:** If Sally runs under **READ UNCOMMITTED**, she could see a price 3.50 even if Joe later aborts



# **Constraints**

## **Ch. 7**

# Constraints and Triggers

- A **constraint** is a relationship among data elements that the DBMS is required to enforce:
  - **Example:** key constraints
- **Triggers** are only executed when a specified condition occurs, e.g., insertion of a tuple:
  - Easier to implement than complex constraints

# Kinds of Constraints

- **Keys**
- **Foreign-key**, or **referential-integrity**
- **Value-based** constraints
  - Constrain values of a particular attribute
- **Tuple-based** constraints
  - Relationship among components
- **Assertions**: any SQL boolean expression

# Review: Single-Attribute Keys

- Place **PRIMARY KEY** or **UNIQUE** after the type in the declaration of the attribute

- **Example:**

```
CREATE TABLE Beers (  
    name CHAR(20) UNIQUE,  
    manf CHAR(20)  
);
```

# Review: Multiattribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         VARCHAR(20) ,  
    price        REAL,  
    PRIMARY KEY (bar, beer)  
);
```

# Foreign Keys

- Values appearing in attributes of one relation must appear together in certain attributes of another relation
- **Example:** in **Sells**(bar, beer, price), we might expect that a **Sells.beer** value also appears in **Beers.name**

# Expressing Foreign Keys

- Use keyword **REFERENCES**, either:
  1. After an attribute (for one-attribute keys).
  2. As an element of the schema:  
**FOREIGN KEY** (<list of attributes>  
**REFERENCES** <relation> (<attributes>)
- Referenced attributes must be declared **PRIMARY KEY** or **UNIQUE**



# Example: With Attribute

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20) REFERENCES Beers(name),  
    price     REAL );
```

# Example: As Schema Element

```
CREATE TABLE Beers (  
    name        CHAR(20) PRIMARY KEY,  
    manf        CHAR(20) );
```

```
CREATE TABLE Sells (  
    bar         CHAR(20),  
    beer        CHAR(20),  
    price       REAL,  
    FOREIGN KEY (beer) REFERENCES Beers (name) );
```

# Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from relation  $R$  (*Foreign/references*), to relation  $S$  (*primary/referenced*) two violations are possible:
  1. An insert or update to  $R$  introduces values not found in  $S$
  2. A deletion or update to  $S$  causes some tuples of  $R$  to “dangle”

# Actions Taken --- (1)

- **Example:** suppose  $R = \text{Sells} \langle \text{bar}, \text{beer}, \text{price} \rangle$ ,  $S = \text{Beers}$ .
- An insert or update to **Sells** that introduces a nonexistent beer must be rejected!
- A deletion or update to **Beers**  $\langle \text{name}, \text{manf} \rangle$  that removes a beer value found in some tuples of **Sells**  $\langle \text{bar}, \text{beer}, \text{price} \rangle$  can be handled in three ways (next slide):

## Actions Taken --- (2)

1. **Default:** Reject the modification
2. **Cascade:** Make the same changes in Sells:
  - ☐ Deleted beer tuple: delete Sells tuple
  - ☐ Updated beer tuple: change value in Sells
3. **Set NULL:** Change the beer attribute in R/Sells tuple to NULL

# Example: Cascade

- Delete the Bud tuple from **Beers**:
  - Then delete all tuples from **Sells** that have beer = 'Bud'
- Update the Bud tuple in **Beers** by changing 'Bud' to 'Budweiser':
  - Then change all **Sells** tuples with beer = 'Bud' to beer = 'Budweiser'

# Example: Set NULL

- Delete the Bud tuple from **Beers**:
  - Change all tuples of **Sells** that have beer = 'Bud' to have beer = NULL
- Update the Bud tuple in **Beers** by changing 'Bud' to 'Budweiser':
  - Same change in **Sells** as we did for deletion

# Choosing a Policy

- When we declare a foreign key, we may choose policies **SET NULL** or **CASCADE** independently for deletions and updates
- Follow the foreign-key declaration by:  
**ON [UPDATE, DELETE][SET NULL or CASCADE]**
- Two such clauses may be used
- Otherwise, the **default** (reject) is used



# Example: Setting Policy

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL,  
    FOREIGN KEY (beer)  
        REFERENCES Beers (name)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE  
);
```

Foreign key

Primary key

# Attribute-Based Checks

- Constraints on the value of a particular attribute
- Add **CHECK**(<condition>) to the declaration for the attribute
- The condition may use the name of the attribute, but any other relation or attribute name must be in a subquery

# Example: Attribute-Based Check

```
CREATE TABLE Sells (  
    bar          CHAR(20),  
    beer         CHAR(20) CHECK ( beer IN  
                                (SELECT name FROM Beers) ),  
    price        REAL      CHECK ( price <= 5.00 )  
);
```

# Timing of Checks

- Attribute-based checks are performed only when a value for that attribute is inserted or updated
  - **Example: CHECK** (`price <= 5.00`) checks every new price and rejects the modification (for that tuple) if the price is more than \$5
  - **Example: CHECK** (`beer IN (SELECT name FROM Beers)`) not checked if a beer is deleted from Beers (unlike foreign-keys)

# Tuple-Based Checks

- **CHECK** (<condition>) may be added as a relation-schema element
- The condition may refer to any attribute of the relation:
  - But other attributes or relations require a subquery
- Checked on **insert** or **update** only

# Example: Tuple-Based Check

- Only Joe's Bar can sell beer **and** for more than \$5 (i.e., check that either [bar is "Joe's Bar"] or [price <= \$5] is TRUE):

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20) ,  
    price        REAL ,  
    CHECK       (bar = 'Joe''s Bar' OR  
                  price <= 5.00)  
);
```

# Assertions

- These are database-schema elements, like relations or views

- Defined by:

**CREATE ASSERTION** <name>

**CHECK** (<condition>);

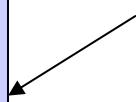
- Condition may refer to any relation or attribute in the database schema

# Example: Assertion

- In **Sells**(bar, beer, price), no bar may charge an average price of more than \$5

```
CREATE ASSERTION NoRipoffBars CHECK (  
  NOT EXISTS (  
    SELECT bar FROM Sells  
    GROUP BY bar  
    HAVING AVG(price) > 5.00  
  ));
```

Bars with an  
average price  
above \$5



- Subquery returns bar(s) with average price > \$5. **Check()** evaluates to TRUE if there is no bar(s) returned from the subquery!



# Example: Assertion

- In **Drinkers**(name, addr, phone) and **Bars**(name, addr, license), there cannot be more bars than drinkers

```
CREATE ASSERTION FewBar CHECK (  
    (SELECT COUNT(*) FROM Bars) <=  
    (SELECT COUNT(*) FROM Drinkers)  
);
```

- Check evaluates to TRUE if # of Bars <= # of Drinkers

# Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database
- A clever system can observe that only certain changes could cause a given assertion to be violated:
  - **Example:** No change to Beers can affect **FewBar!** Neither can an insertion to Drinkers can change the Assertion to be False



# Triggers

## Ch. 7

# Triggers: Motivation

- Assertions are powerful, but the DBMS often can't tell when they need to be checked
- Attribute- and tuple-based checks are checked at known times, but are not powerful
- Triggers let the user decide when to check for any condition

# Event-Condition-Action Rules

- Another name for “trigger” is *ECA rule*, or *Event-Condition-Action rule*
- *Event*: typically a type of database modification, e.g., “insert on Sells”
- *Condition*: Any SQL boolean-valued expression
- *Action*: Any SQL statements

# Preliminary Example: A Trigger

- Instead of using a foreign-key constraint and rejecting insertions into **Sells**(bar, beer, price)  $\leftarrow$  (*Foreign/references*) with unknown beers (*primary/referenced*), a trigger can add that beer to Beers, with a NULL manufacturer 😊

# Example: Trigger Definition

**CREATE TRIGGER** BeerTrig

**AFTER INSERT ON Sells**

The **E**vent

**REFERENCING NEW ROW AS NewTuple**  
**FOR EACH ROW**

**WHEN** (NewTuple.beer **NOT IN**  
**(SELECT name FROM Beers))**

The **C**ondition

**INSERT INTO Beers(name)**  
**VALUES(NewTuple.beer);**

The **A**ction

< If inserting new Beer in the “Sells” table that is not known in the “Beers” table, then insert tuple with the new Beer in the “Beers” Table >

# Options: CREATE TRIGGER (1)

**CREATE TRIGGER** <name>

■ Or:

**CREATE OR REPLACE TRIGGER** <name>

- Useful if there is a trigger with that name and you want to modify the existing trigger



# Options: The Event (2)

- **AFTER** can be **BEFORE**

- Also, **INSTEAD OF**, if the relation is a view
  - A clever way to execute view modifications: have triggers translate them to appropriate modifications on the base tables

- **INSERT** can be **DELETE** or **UPDATE**

- And **UPDATE** can be **UPDATE ... ON** a particular attribute

## Options: REFERENCING (3)

- **INSERT** statements imply a new tuple (for row-level) or new table (for statement-level):
  - The “table” is the set of inserted tuples
- **DELETE** implies an old tuple or table
- **UPDATE** implies both, i.e., Delete and Insert
- Refer to these by  
[NEW OLD][TUPLE TABLE] **AS** <name>

# Options: FOR EACH ROW (4)

- Triggers (E-C-A) are either “row-level” or “statement-level”:
  - **FOR EACH ROW** indicates row-level; its absence indicates statement-level
  - **Row level triggers:** execute once for each modified tuple
  - **Statement-level triggers:** execute once for a SQL statement, regardless of how many tuples are modified

# Options: The Condition (5)

- Any boolean-valued condition
- Evaluated on the database as it would exist before or after the triggering event, depending on whether **BEFORE** or **AFTER** is used
  - But always before the changes take effect
- Access the new/old tuple/table through the names in the **REFERENCING** clause

# Options: The Action (6)

- There can be more than one SQL statement in the action:
  - Surround by **BEGIN ... END** if there is more than one
- But **queries** make no sense in an action, so we are really limited to modifications

# Another Example

- Using **Sells**(bar, beer, price) and a unary relation **RipoffBars**(bar), maintain a list of bars that raise the price of any beer by more than \$1

# The Trigger

**CREATE TRIGGER** PriceTrig

The event –  
only changes  
to prices

**AFTER UPDATE** OF price **ON Sells**

**REFERENCING**  
**OLD ROW AS** ooo  
**NEW ROW AS** nnn

Updates let us  
talk about old  
and new tuples

**FOR EACH ROW**

We need to consider  
each price change

**WHEN**(nnn.price > ooo.price + 1.00)

Condition:  
a raise in  
price > \$1

**INSERT INTO** RipoffBars  
**VALUES**(nnn.bar);

When the price change  
is great enough, add  
the bar to RipoffBars



# **Views**

## **Ch. 8**



# Views

- A **view** is a relation defined in terms of stored tables (called **base tables**) and other views
- **Two kinds:**
  1. **Virtual** = not stored in the database; just a query for constructing the relation
  2. **Materialized** = actually constructed and stored

# Declaring Views

- Declare by:

```
CREATE [MATERIALIZED] VIEW  
    <name> AS <query>;
```

- Default is virtual

# Example: View Definition

- **CanDrink**(drinker, beer) is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

```
CREATE VIEW CanDrink AS
  SELECT drinker, beer
  FROM Frequents, Sells
  WHERE Frequents.bar = Sells.bar;
```

# Example: Accessing a View

- Query a view as if it were a base table:
  - **Also:** a limited ability to modify views if it makes sense as a modification of one underlying base table
- Example query:

```
SELECT beer FROM CanDrink
WHERE drinker = 'Sally';
```

# Triggers on Views

- Generally, it is impossible to modify a virtual view, because it doesn't exist
- But an **INSTEAD OF** trigger lets us interpret view modifications in a way that makes sense
- **Example: View Synergy** has (**drinker**, **beer**, **bar**) triples such that the bar serves the beer, the drinker frequents the bar, and drinker likes the beer

# Example: The View

**CREATE VIEW** Synergy **AS**

Pick one copy of  
each attribute

**SELECT** Likes.drinker, Likes.beer, Sells.bar

**FROM** Likes, Sells, Frequents

**WHERE** Likes.drinker = Frequents.drinker

**AND** Likes.beer = Sells.beer

**AND** Sells.bar = Frequents.bar;

Natural join of Likes,  
Sells, and Frequents

# Interpreting a View Insertion

- We cannot insert into Synergy --- it is a virtual view
- But we can use an **INSTEAD OF** trigger to turn a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents:
  - Sells.price will have to be NULL.

# The Trigger

```
CREATE TRIGGER ViewTrig
  INSTEAD OF INSERT ON Synergy
  REFERENCING NEW ROW AS n
  FOR EACH ROW
  BEGIN
    INSERT INTO LIKES VALUES(n.drinker, n.beer);
    INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
    INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
  END;
```



# Materialized Views

- **Problem:** each time a base table changes, the materialized view may change:
  - Cannot afford to re-compute the view with each change
- **Solution:** Periodic reconstruction of the materialized view, which is otherwise “out of date”

# Example: Axxess/Class Mailing List

- The class mailing list **cs510-aut0708-students** is in effect a materialized view of the class enrollment in Axxess
- **Actually updated four times/day:**
  - You can enroll and miss an email sent out after you enroll

# Example: A Data Warehouse

- Wal-Mart stores every sale at every store in a database
- Overnight, the sales for the day are used to update a *data warehouse* = materialized views of the sales
- The warehouse is used by analysts to predict trends and move goods to where they are selling best



# **Indexes**

## **Ch. 8**

# Indexes

- **Index** = data structure used to speed access to tuples of a relation, given values of one or more attributes
- Could be a hash table, but in a DBMS it is always a balanced search tree with giant nodes (a full disk page) called a **B-tree**

# Declaring Indexes

- No standard!

- Typical syntax:

```
CREATE INDEX BeerInd ON  
           Beers (manf) ;
```

```
CREATE INDEX SellInd ON  
           Sells (bar, beer) ;
```

# Using Indexes

- Given a value  $v$ , the index takes us to only those tuples that have  $v$  in the attribute(s) of the index
- **Example:** use **BeerInd** (index on attribute ( $manf$ ) in **Beers** table) and **SellInd** {index on attribute ( $bar$ ,  $beer$ ) in **Sells** table} to find the prices of beers manufactured by Pete's and sold by Joe's bar. → next slide

# Using Indexes --- (2)

```
SELECT price
FROM   Beers, Sells
WHERE  manf = 'Pete''s' AND
       Beers.name = Sells.beer AND
       bar = 'Joe''s Bar';
```

1. Use **BeerInd** to get all the beers made by Pete's
2. Then use **SellInd** to get prices of those beers (i.e., **<beer, price>**), with bar = 'Joe's Bar'
3. Join output of steps (1) and (2) to get the price of beers that are made by Pete's and are sold in Joe's Bar



# Database Tuning

- A major problem in making a database run fast is deciding which indexes to create?
- **Pro:** An index speeds up queries that can use it
- **Con:** An index slows down all modifications on its relation because the index must be modified too

# Example: Tuning

- Suppose the only things we did with our beers database was:
  1. Insert new facts into a relation (10%)
  2. Find the price of a given beer at a given bar (90%)
- Then **SellInd** on **Sells**(bar, beer) → {READ} would be wonderful, but **BeerInd** on **Beers**(manf) → {INSERT} would be harmful

# Tuning (Design) Advisors

- A major research thrust:
  - Because hand tuning is so hard
- An advisor gets a *query load*, e.g.:
  1. Choose random queries from the history of queries run on the database, or
  2. Designer provides a sample workload

# Tuning Advisors --- (2)

- The “**design advisor**” generates candidate indexes and evaluates each on the workload:
  - Feed each sample query to the query optimizer, which assumes only this one index is available
  - Measure the improvement/degradation in the average running time of the queries



**END**